

Divyansh Khanna

divyanshk.github.io
dk3399@nyu.edu | 646-201-6265 | New York, NY

EDUCATION

NEW YORK UNIVERSITY

MASTERS IN COMPUTER SCIENCE
COURANT INSTITUTE OF
MATHEMATICAL SCIENCES
Expected May 2019 | New York City
GPA: 3.83

BITS PILANI

BACHELOR OF ENGINEERING IN
COMPUTER SCIENCE
MASTER OF SCIENCE IN
MATHEMATICS
July 2016 | Goa, India
Cum. GPA: 8.17 / 10.0
CS Major GPA: 8.80 / 10.0

SKILLS

PROGRAMMING

Java • Python • C++

Knowledge of:

Tensorflow • Scikit-learn

PyTorch • Hive

Cassandra • Elasticsearch

Familiar with:

Spark • Scala

COURSEWORK

GRADUATE

Deep Generative Models
HPC for Machine Learning
Fundamental Algorithms
Computer Vision

UNDERGRADUATE

Design and Analysis of Algorithms
Artificial Intelligence
(Teaching Assistant)
Parallel Computing
Optimization
Graphs and Networks

LINKS

Github:// [divyanshk](#)

LinkedIn:// [divyanshkhanna](#)

Twitter:// [@divyanshkhanna1](#)

EXPERIENCE

FLIPKART | SOFTWARE DEVELOPMENT ENGINEER

December 2016 – July 2017 | Bengaluru, India

- Worked with the Data Platform team on building products for better insights on daily data captured across the Flipkart ecosystem
- Developed new features for the Common Data Model, a framework to provide a self sustaining platform capturing complete data life cycles
- Identified the bottleneck of a production Hive job, designed and implemented a solution to cut the CPU time by over 5 times

INDIAN INSTITUTE OF SCIENCE | PROJECT ASSISTANT

July 2016 – December 2016 | Bengaluru, India

- Collaborated with the Energy Analytics team on behavioral activity models
- Built models for analyzing households' power consumption activities
- Deployed ensemble supervised learning along with hidden Markov models to model the consumption patterns

MYNTRA | SOFTWARE ENGINEERING INTERN

Jan 2016 – July 2016 | Bengaluru, India

- Developed a dashboard for detailed topic analysis of the Myntra mobile app's personalized feed using NoSQL databases and JS backend framework
- Created a tool for fetching the top posts within a date range ranked by various user provided metrics
- The dashboard was used across multiple product and data science teams to validate the feed's performance

SIEMENS RESEARCH | SUMMER INTERN

May 2015 – July 2015 | Bengaluru, India

- Identified performance overheads of CUDA and cpp backend of Thrust: A template library for CUDA
- Used timing and profiling tools to estimate the performance differences
- The analysis was used to overcome delays in production code

PROJECTS

ADVERSARIAL BEHAVIOR CLONING

- Course project for CSCI-GA.3033-022 Deep Generative Models
- Developed a adversarial auto-encoder based behavior cloning technique to imitate expert demonstrations

PARALLEL APPROACH TO SEARCH ALGORITHMS FOR DISCRETE OPTIMIZATION PROBLEMS

- Worked with parallel search algorithms namely Depth First, Best First, Iterative Deepening A Star and their application to optimization tasks using OpenMPI
- Covered load balancing and communication analysis of the parallel implementations of the algorithms in Python